

FULL SEO AUDIT REPORT

aplory.dev

Prepared by Claude SEO

58/100

SEO Health Score

August 28, 2026 Full Audit

Table of Contents

1. Executive Summary

Key Metrics, Critical Issues & Quick Wins

2. Audit Categories

What Works

Findings by Severity

3. Action Plan

Phased Roadmap

4. Data Sources & Methodology

1. Executive Summary

This report presents a comprehensive SEO audit of **aplory.dev** for a **SaaS (B2B automated-response platform) with 9 industry-vertical landing pages, targeting small businesses in Serbia** site, generated on August 28, 2026. Findings combine technical, content, schema, performance, visual, and search-readiness evidence as available.

58/100

SEO HEALTH SCORE

Critical Issues Found:

1. **Info:** Pure client-side rendering: homepage + 10 niche pages ship an empty `<div id="root"></div>` with no server-rendered content, confirmed via curl and in the dist/ build.
2. **Info:** Product-claim vs. search-intent mismatch: niche pages imply live call-answering, but the homepage FAQ admits the voice agent doesn't exist yet; Serbian SERP is dominated by voice-agent competitors.
3. **Info:** Identical H2 duplicated across 10 URLs; vertical pages carry only ~150-200 unique words each on a shared template.
4. **Info:** Zero trust signals on any vertical page: no FAQ, testimonials, pricing, client logos, About page, or company registration number.
5. **Info:** AI Search Readiness (GEO) scores 34/100, driven almost entirely by the CSR issue plus absence of any citable-length (134-167 word) passages.

Quick Wins:

1. Defer/lazy-load the LeadConnector chat widget script (largest render-blocking resource, ~1,340ms)
2. Publish /llms.txt with a plain-text summary of APLORY
3. Compress logo-mark.png (19.6KB of 20.9KB is waste)
4. Add missing security headers via netlify.toml [[headers]]
5. Reformat the '62% of missed calls' stat into one fully attributed, disclosed sentence

2. Audit Categories

2.1 Technical SEO

Score: **61/100**

What Works

- robots.txt cleanly allows all UAs and correctly declares sitemap.xml
- sitemap.xml valid, 15 URLs, sensible priority weighting and lastmod
- Every page has correct self-referential canonical, unique title/meta description
- Full OG/Twitter tags present as static tags (unaffected by CSR issue)
- Clean single-hop HTTPS/www redirects, hyphenated URL structure
- Sound cache-control strategy on Netlify; JSON-LD ProfessionalService schema server-rendered on homepage

Findings

Homepage and all 10 niche landing pages ship empty `<div id="root"></div>` - pure CSR, no SSR/prerendering **Critical**

Confirmed via curl and in both repo source and dist/ build: 11 of 15 sitemap URLs (all commercially important pages) have zero content in raw HTML. Only the 4 legal pages are genuinely static. No SSG/prerendering plugin configured in vite.config.js despite each route already having its own HTML entry point.

Recommendation: Add prerendering/SSG (vite-plugin-ssg, vite-plugin-prerender-html, or a custom build-time render step) so each dist/*.html ships fully rendered markup; React can still hydrate on top.

No custom security headers beyond default HSTS **High**

No Content-Security-Policy, X-Content-Type-Options, X-Frame-Options, Referrer-Policy, or Permissions-Policy. HSTS present but missing includeSubDomains/preload.

Recommendation: Add a netlify.toml [[headers]] block or public/_headers file applied to /* with the missing headers, scoped to allow existing third-party origins (leadconnectorhq.com, fonts.googleapis.com).

LCP gated by JS execution, compounding the CSR issue **Medium**

Hero images are correctly dimensioned but can't paint until the JS bundle downloads/parses/executes/hydrates, since they only exist in the DOM post-render.

Recommendation: Fixing the prerendering issue resolves most of this; add `fetchpriority="high"` to the hero image once part of initial render.

IndexNow protocol not implemented **Medium**

No IndexNow key file and no post-deploy submission to api.indexnow.org.

Recommendation: Add a key file plus a post-deploy hook submitting the 15 sitemap URLs after each Netlify deploy, to speed Bing/Yandex re-crawl.

Inconsistent meta name="robots" declaration **Low**

Only the 4 legal pages explicitly declare meta robots index,follow; no functional difference from the default.

Recommendation: Cosmetic only - standardize either way, no priority.

2.2 Content Quality

Score: **42/100**

What Works

- Readability is good - short sentences, concrete scenarios, plain Serbian language
- Demo chat transcripts are genuinely vertical-specific and a real trust/experience signal
- Homepage FAQ candidly admits the voice-agent feature isn't ready yet - a rare, credible trust signal
- 62% statistic is at least cited with a working source URL, better than an unsourced claim

Findings

No FAQ/Q&A content on any of the 9 vertical pages **Critical**

The Faq component only renders on the homepage. Vertical pages have zero structured, quotable Q&A content, the highest-value format for featured snippets and LLM citation.

Recommendation: Add 3-5 vertical-specific FAQ entries per page as prose content (pricing model, integrations, compliance question relevant to that industry).

Exact duplicate H2 across 10 pages **Critical**

NichePain.jsx hardcodes the identical sentence as the section heading on all 9 niche pages, and the homepage uses the same sentence for problem.title.

Recommendation: Write a distinct, vertical-specific H2 for each niche page using the existing pain-card copy as raw material.

No About/company/team page or entity trust signals **High**

Only a personal Gmail address and a mobile number are exposed. No founder identity, headcount, testimonials, or third-party recognition anywhere.

Recommendation: Add an About/Company page with founder details, physical address (currently buried only in legal pages), and PIB/company registration number.

Core statistic is old, small-sample, market-mismatched, and undisclosed **High**

The 62% headline stat is sourced to a 2016, 85-company U.S. study, presented without disclosure that it isn't Serbian or current data.

Recommendation: Either caveat the stat explicitly as a U.S. benchmark, commission Serbia-relevant data, or replace with the company's own aggregated client data.

"How it works" explainer 100% identical across all 9 vertical pages and homepage **Medium**

NicheHow.jsx pulls identical eyebrow, title, and step copy from shared sr.js data - nothing niche-specific except the demo chat transcript.

Recommendation: Localize at least the three step bodies per vertical.

No FAQPage schema despite homepage FAQ content **Medium**

Homepage FAQ has 4 good Q&As but no structured data wrapping it.

Recommendation: Add FAQPage schema as a content-organization aid only - do not expect a Google SERP benefit, as FAQ rich results were retired May 2026.

Thin overall word count relative to search intent **Medium**

Genuinely unique body copy per vertical page is ~150-200 words; sta-radimo.html largely restates homepage copy rather than going deeper.

Recommendation: Expand sta-radimo.html with mechanism, SLA, and data-handling depth.

2.3 On-Page SEO

Score: 58/100

What Works

- Unique titles and meta descriptions per page, appropriately localized
- Clean, descriptive, hyphenated URL structure with no query strings
- Correct self-referential canonicals across all 15 pages

Findings

Identical H2 and heading structure across 10 URLs **High**

Shared template produces near-identical on-page heading hierarchy across all vertical pages, undermining topical differentiation signals.

Recommendation: See Content Quality findings - de-duplicate H2s and localize section copy per vertical.

No internal linking depth beyond primary navigation **Medium**

No contextual cross-links between related vertical pages or from blog/resource content (none exists) into the commercial pages.

Recommendation: Once content depth increases (FAQ, expanded sta-radimo.html), add contextual internal links between related pages.

2.4 Schema / Structured Data

Score: 72/100

What Works

- Server-rendered ProfessionalService + OfferCatalog/Service/Offer JSON-LD, valid, no deprecated types
- Per-niche-page Service schema referencing the same provider with areaServed: RS

Findings

Missing image/logo, sameAs, and address properties **Info**

ProfessionalService schema lacks image, sameAs (social/GBP links), and a PostalAddress, all of which support Knowledge Panel eligibility and entity consolidation.

Recommendation: Add image (reuse existing og:image), sameAs with real profile URLs, and address with at least addressCountry: RS.

No FAQPage, Review/AggregateRating, or LocalBusiness/GeoCoordinates schema **Medium**

Despite areaServed: RS and the product serving inherently local-service verticals, no LocalBusiness-family schema or review markup exists.

Recommendation: Add once testimonials/reviews exist; add LocalBusiness/geo schema given the explicit RS service area.

2.5 Performance

Score: 94/100

What Works

- LCP 2.0-2.4s (Pass), CLS 0.03-0.07 (Pass), TBT 20-50ms (Pass) across all 3 pages tested, lab data
- TTFB excellent (67-184ms) with Brotli and Netlify edge caching confirmed working
- No legacy/duplicated JavaScript shipped

Findings

Third-party chat widget script is the largest render-blocking resource **High**

widgets.leadconnectorhq.com/loader.js costs ~1,340ms of blocking time; nearly all LCP time (2,086/2,154ms on homepage) sits in element-render-delay rather than TTFB.

Recommendation: Load the chat widget with defer or on user interaction/scroll rather than blocking initial render.

Duplicate font loading (Google Fonts + Bunny Fonts) both render-blocking **High**

Google Fonts CSS costs ~1,418ms; a separate Bunny Fonts Roboto file costs ~685ms - two font origins loading synchronously ahead of render.

Recommendation: Consolidate to one font provider, ideally self-hosted given the existing self-hosted AdwaitaSans.woff2 precedent.

Unoptimized logo-mark.png **Medium**

19.6KB of the file's 20.9KB is wasted - 2x oversized PNG, not in a modern format.

Recommendation: Re-export as properly sized WebP/AVIF via the existing Vite asset pipeline.

~17KB unused CSS per page **Medium**

Shared base CSS bundle ships more rules than any single route needs, consistent across all 3 pages tested.

Recommendation: Verify Tailwind content purge globs cover all 15 HTML entry points; consider per-route CSS chunking.

Short cache lifetimes on third-party chat-widget assets **Medium**

8 third-party resources cache for only 4 hours, ~74KB re-download waste per repeat-visit cache miss.

Recommendation: Third-party controlled; reinforces the case for deferring/lazy-loading the widget instead.

2.6 AI Search Readiness (GEO)

Score: **34/100**

What Works

- robots.txt cleanly allows all AI crawlers (GPTBot, ClaudeBot, PerplexityBot, Google-Extended, etc.)
- sitemap.xml current and correctly referenced
- Server-rendered ProfessionalService JSON-LD gives non-JS crawlers a usable baseline summary

Findings

Homepage and all niche pages are fully client-side rendered with an empty body as served **Critical**

OpenAI's, Anthropic's, and Perplexity's crawlers have little-to-no reliable JS execution, so the page is effectively text-empty on fetch for those platforms. This caps every other GEO dimension.

Recommendation: Statically pre-render each route - same fix as the Technical SEO Critical finding.

No /llms.txt **High**

GET <https://aplory.dev/llms.txt> returns 404. Optional and doesn't affect Google, but is one of the few channels to hand non-JS crawlers a clean summary today.

Recommendation: Publish a plain-text llms.txt summarizing what APLORY is, who it serves, and key page links.

No long-form, self-contained citable passages anywhere **High**

Optimal AI-citation passage length is ~134-167 words; every extractable answer block on the homepage falls well short (16-36 words).

Recommendation: Add 1-2 expanded 134-167 word paragraphs as prose content, not as FAQPage schema (no Google ranking benefit from the markup itself post-May-2026 retirement).

Zero external brand/entity signals **High**

No Wikipedia entity, Reddit mentions, YouTube presence, or LinkedIn company page found. YouTube/Reddit presence correlates more strongly with AI citation than domain authority.

Recommendation: Build minimal external entity presence over time: demo/explainer video, LinkedIn company page, encourage early-customer mentions.

2.7 Images

Score: **70/100**

What Works

- Hero images correctly sized with explicit width/height attributes (good CLS practice)
- loading="eager" correctly applied to above-the-fold hero images

Findings

logo-mark.png oversized and wrong format **Medium**

400x81px PNG displayed at 196x40px (2x oversized), not in WebP/AVIF - 95% of its bytes are unnecessary.

Recommendation: Re-export as WebP/AVIF at correct dimensions.

3. Action Plan

Phase 1: Critical Fixes Week 1

- Add prerendering/SSG for the 11 CSR routes so real content ships in initial HTML
- Resolve the voice-vs-text product-claim mismatch on all 9 vertical pages
- De-duplicate the identical H2 across all 10 pages with vertical-specific headlines

Phase 2: High-Impact Improvements Weeks 2-3

- Defer/lazy-load the LeadConnector chat widget script
- Consolidate font loading to one provider
- Add missing security headers via netlify.toml
- Add an About/Company section with founder, address, and PIB
- Add 3-5 vertical-specific FAQ entries per niche page (prose, not schema)
- Reformat and disclose the '62% of missed calls' statistic
- Publish /llms.txt

Phase 3: Content & Authority Month 2

- Localize 'How it works' step copy per vertical
- Expand sta-radimo.html with mechanism/SLA/data-handling depth
- Add directional trust proof (client count or quotes per vertical)
- Publish minimal pricing/packaging
- Add expanded 134-167 word citable paragraphs
- Compress logo-mark.png to WebP/AVIF
- Purge unused CSS (~17KB/page)
- Begin Serbian-market link-building (Startit.rs, Product Hunt, PKS directories)
- Add Organization schema with sameAs, image, and address

Phase 4: Monitoring & Iteration Ongoing

- Connect Google Search Console and a PageSpeed Insights API key
- Add a free Moz API key to move backlink tracking to Tier 1
- Verify mobile hamburger menu tap-target size in a real browser

- Re-run the full audit after Phase 1-2 land
- Set an SEO drift baseline

Data Sources & Methodology

Source	Description	Update Frequency
PageSpeed Insights API	Lighthouse lab audit (mobile emulation, Moto G Power, slow 4G)	Real-time
Chrome UX Report (CrUX)	28-day rolling field data from real Chrome users	Daily ~04:00 UTC
CrUX History API	25-week p75 trend data per metric	Weekly
Google Search Console	Search Analytics (clicks, impressions, CTR, position)	2-3 day lag
URL Inspection API	Per-URL index status, coverage state, crawl info	Real-time (2,000/day)

Report generated by Claude SEO — Google SEO Intelligence Skill — August 28, 2026
Methodology based on Google Web Vitals thresholds, Search Console documentation, and Lighthouse scoring algorithms.